cm010 Exercises: Tibble Joins

Requirements

You will need Joey's `singer` R package for this exercise. And to install that, you'll need to install `devtools`. Running this code in your console should do the trick:

```
install.packages("devtools")
devtools::install_github("JoeyBernhardt/singer")
```

Load required packages:

Exercise 1: `singer`

The package `singer` comes with two smallish data frames about songs. Let's take a look at them (after minor modifications by renaming and shuffling):

```
(time <- as_tibble(songs) %>%
   rename(song = title))
```

```
## # A tibble: 22 x 3
##    song                              artist_name      year
##    <chr>                             <chr>           <int>
##  1 Corduroy                          Pearl Jam        1994
##  2 Grievance                         Pearl Jam        2000
##  3 Stupidmop                         Pearl Jam        1994
##  4 Present Tense                     Pearl Jam        1996
##  5 MFC                               Pearl Jam        1998
##  6 Lukin                             Pearl Jam        1996
##  7 It's Lulu                         The Boo Radleys  1995
##  8 Sparrow                           The Boo Radleys  1992
##  9 Martin_ Doom! It's Seven O'Clock The Boo Radleys  1995
## 10 Leaves And Sand                   The Boo Radleys  1993
## # ... with 12 more rows
```

```
(album <- as_tibble(locations) %>%
   select(title, everything()) %>%
   rename(album = release,
          song  = title))
```

```
## # A tibble: 14 x 4
##    song                        artist_name   city        album
##    <chr>                       <chr>         <chr>       <chr>
##  1 Grievance                   Pearl Jam     Seattle, WA Binaural
##  2 Stupidmop                   Pearl Jam     Seattle, WA Vitalogy
##  3 Present Tense               Pearl Jam     Seattle, WA No Code
##  4 MFC                         Pearl Jam     Seattle, WA Live On Two Legs
##  5 Lukin                       Pearl Jam     Seattle, WA Seattle Washingto~
##  6 Stuck On Amber              The Boo Radl~ Liverpool,~ Wake Up!
##  7 It's Lulu                   The Boo Radl~ Liverpool,~ Best Of
##  8 Sparrow                     The Boo Radl~ Liverpool,~ Everything's Alri~
##  9 High as Monkeys             The Boo Radl~ Liverpool,~ Kingsize
## 10 Butterfly McQueen           The Boo Radl~ Liverpool,~ Giant Steps
## 11 My One and Only Love        Carly Simon   New York, ~ Moonlight Serenade
## 12 It Was So Easy  (LP Versio~ Carly Simon   New York, ~ No Secrets
```

```
## 13 I've Got A Crush On You     Carly Simon   New York, ~ Clouds In My Coff~
## 14 "Manha De Carnaval (Theme ~ Carly Simon   New York, ~ Into White
```

1. We really care about the songs in `time`. But, which of those songs do we know its corresponding album? To not show the coreeponsding albums:

```
time %>%
  semi_join(album, by = c("song", "artist_name"))
```

```
## # A tibble: 13 x 3
##    song                                                artist_name      year
##    <chr>                                               <chr>           <int>
##  1 Grievance                                           Pearl Jam        2000
##  2 Stupidmop                                           Pearl Jam        1994
##  3 Present Tense                                       Pearl Jam        1996
##  4 MFC                                                 Pearl Jam        1998
##  5 Lukin                                               Pearl Jam        1996
##  6 It's Lulu                                           The Boo Radleys  1995
##  7 Sparrow                                             The Boo Radleys  1992
##  8 High as Monkeys                                     The Boo Radleys  1998
##  9 Butterfly McQueen                                   The Boo Radleys  1993
## 10 My One and Only Love                                Carly Simon      2005
## 11 It Was So Easy  (LP Version)                        Carly Simon      1972
## 12 I've Got A Crush On You                             Carly Simon      1994
## 13 "Manha De Carnaval (Theme from \"Black Orpheus\")" Carly Simon      2007
```

To show all corresponding albums

```
time %>%
  inner_join(album, by = c("song","artist_name"))
```

```
## # A tibble: 13 x 5
##    song                    artist_name    year city       album
##    <chr>                   <chr>         <int> <chr>      <chr>
##  1 Grievance               Pearl Jam      2000 Seattle, ~ Binaural
##  2 Stupidmop               Pearl Jam      1994 Seattle, ~ Vitalogy
##  3 Present Tense           Pearl Jam      1996 Seattle, ~ No Code
##  4 MFC                     Pearl Jam      1998 Seattle, ~ Live On Two Legs
##  5 Lukin                   Pearl Jam      1996 Seattle, ~ Seattle Washing~
##  6 It's Lulu               The Boo Radl~  1995 Liverpool~ Best Of
##  7 Sparrow                 The Boo Radl~  1992 Liverpool~ Everything's Al~
##  8 High as Monkeys         The Boo Radl~  1998 Liverpool~ Kingsize
##  9 Butterfly McQueen       The Boo Radl~  1993 Liverpool~ Giant Steps
## 10 My One and Only Love    Carly Simon    2005 New York,~ Moonlight Seren~
## 11 It Was So Easy  (LP Ver~ Carly Simon    1972 New York,~ No Secrets
## 12 I've Got A Crush On You  Carly Simon    1994 New York,~ Clouds In My Co~
## 13 "Manha De Carnaval (The~ Carly Simon    2007 New York,~ Into White
```

2. Go ahead and add the corresponding albums to the `time` tibble, being sure to preserve rows even if album info is not readily available.

```
time %>%
  left_join(album, by = c("song","artist_name"))
```

```
## # A tibble: 22 x 5
##    song                artist_name     year city        album
##    <chr>               <chr>          <int> <chr>       <chr>
##  1 Corduroy            Pearl Jam       1994 <NA>        <NA>
```

```
##  2 Grievance           Pearl Jam       2000 Seattle, WA Binaural
##  3 Stupidmop           Pearl Jam       1994 Seattle, WA Vitalogy
##  4 Present Tense       Pearl Jam       1996 Seattle, WA No Code
##  5 MFC                 Pearl Jam       1998 Seattle, WA Live On Two Legs
##  6 Lukin               Pearl Jam       1996 Seattle, WA Seattle Washington~
##  7 It's Lulu           The Boo Radle~  1995 Liverpool,~ Best Of
##  8 Sparrow             The Boo Radle~  1992 Liverpool,~ Everything's Alrig~
##  9 Martin_ Doom! It's~ The Boo Radle~  1995 <NA>        <NA>
## 10 Leaves And Sand     The Boo Radle~  1993 <NA>        <NA>
## # ... with 12 more rows
```

3. Which songs do we have "year", but not album info?

```
time %>%
  semi_join(album, by = "song")
```

```
## # A tibble: 13 x 3
##    song                                                artist_name      year
##    <chr>                                               <chr>           <int>
##  1 Grievance                                           Pearl Jam        2000
##  2 Stupidmop                                           Pearl Jam        1994
##  3 Present Tense                                       Pearl Jam        1996
##  4 MFC                                                 Pearl Jam        1998
##  5 Lukin                                               Pearl Jam        1996
##  6 It's Lulu                                           The Boo Radleys  1995
##  7 Sparrow                                             The Boo Radleys  1992
##  8 High as Monkeys                                     The Boo Radleys  1998
##  9 Butterfly McQueen                                   The Boo Radleys  1993
## 10 My One and Only Love                                Carly Simon      2005
## 11 It Was So Easy  (LP Version)                        Carly Simon      1972
## 12 I've Got A Crush On You                             Carly Simon      1994
## 13 "Manha De Carnaval (Theme from \"Black Orpheus\")" Carly Simon      2007
```

4. Which artists are in `time`, but not in `album`?

```
time %>%
  anti_join(album, by = "artist_name")
```

```
## # A tibble: 5 x 3
##   song                  artist_name   year
##   <chr>                 <chr>        <int>
## 1 Mine Again            Mariah Carey  2005
## 2 Don't Forget About Us Mariah Carey  2005
## 3 Babydoll              Mariah Carey  1997
## 4 Don't Forget About Us Mariah Carey  2005
## 5 Vision Of Love        Mariah Carey  1990
```

5. You've come across these two tibbles, and just wish all the info was available in one tibble. What would you do?

```
time %>%
  full_join(album, by = c("song", "artist_name"))
```

```
## # A tibble: 23 x 5
##    song                artist_name     year city        album
##    <chr>               <chr>          <int> <chr>       <chr>
##  1 Corduroy            Pearl Jam       1994 <NA>        <NA>
##  2 Grievance           Pearl Jam       2000 Seattle, WA Binaural
```

```
##  3 Stupidmop           Pearl Jam       1994 Seattle, WA Vitalogy
##  4 Present Tense       Pearl Jam       1996 Seattle, WA No Code
##  5 MFC                 Pearl Jam       1998 Seattle, WA Live On Two Legs
##  6 Lukin               Pearl Jam       1996 Seattle, WA Seattle Washington~
##  7 It's Lulu           The Boo Radle~  1995 Liverpool,~ Best Of
##  8 Sparrow             The Boo Radle~  1992 Liverpool,~ Everything's Alrig~
##  9 Martin_ Doom! It's~ The Boo Radle~  1995 <NA>        <NA>
## 10 Leaves And Sand     The Boo Radle~  1993 <NA>        <NA>
## # ... with 13 more rows
```

Exercise 2: LOTR

Load in the three Lord of the Rings tibbles that we saw last time:

```
fell <- read_csv("https://raw.githubusercontent.com/jennybc/lotr-tidy/master/data/The_Fellowship_Of_The
```

```
## Parsed with column specification:
## cols(
##   Film = col_character(),
##   Race = col_character(),
##   Female = col_double(),
##   Male = col_double()
## )
```

```
ttow <- read_csv("https://raz.githubusercontent.com/jennybc/lotr-tidy/master/data/The_Two_Towers.csv")
```

```
## Parsed with column specification:
## cols(
##   Film = col_character(),
##   Race = col_character(),
##   Female = col_double(),
##   Male = col_double()
## )
```

```
retk <- read_csv("https://raw.githubusercontent.com/jennybc/lotr-tidy/master/data/The_Return_Of_The_King
```

```
## Parsed with column specification:
## cols(
##   Film = col_character(),
##   Race = col_character(),
##   Female = col_double(),
##   Male = col_double()
## )
```

1. Combine these into a single tibble.

```
bind_rows(fell, ttow,retk)
```

```
## # A tibble: 9 x 4
##   Film                      Race   Female  Male
##   <chr>                     <chr>   <dbl> <dbl>
## 1 The Fellowship Of The Ring Elf      1229   971
## 2 The Fellowship Of The Ring Hobbit     14  3644
## 3 The Fellowship Of The Ring Man         0  1995
## 4 The Two Towers             Elf       331   513
## 5 The Two Towers             Hobbit      0  2463
## 6 The Two Towers             Man       401  3589
## 7 The Return Of The King     Elf       183   510
```

```
## 8 The Return Of The King     Hobbit      2  2673
## 9 The Return Of The King     Man       268  2459
```

2. Which races are present in "The Fellowship of the Ring" (`fell`), but not in any of the other ones?

```
fell %>%
  anti_join(ttow, by = "Race") %>%
  anti_join(retk, by = "Race")
```

```
## # A tibble: 0 x 4
## # ... with 4 variables: Film <chr>, Race <chr>, Female <dbl>, Male <dbl>
```

Exercise 3: Set Operations

Let's use three set functions: `intersect`, `union` and `setdiff`. We'll work with two toy tibbles named `y` and `z`, similar to Data Wrangling Cheatsheet

```
(y <-  tibble(x1 = LETTERS[1:3], x2 = 1:3))
```

```
## # A tibble: 3 x 2
##   x1       x2
##   <chr> <int>
## 1 A         1
## 2 B         2
## 3 C         3
```

```
(z <- tibble(x1 = c("B", "C", "D"), x2 = 2:4))
```

```
## # A tibble: 3 x 2
##   x1       x2
##   <chr> <int>
## 1 B         2
## 2 C         3
## 3 D         4
```

1. Rows that appear in both `y` and `z` intersect - gets rows that appear in both

```
intersect(y, z)
```

```
## # A tibble: 2 x 2
##   x1       x2
##   <chr> <int>
## 1 B         2
## 2 C         3
```

2. You collected the data in `y` on Day 1, and `z` in Day 2. Make a data set to reflect that. bind rows or full join works here

```
bind_rows(
  mutate(y, day = "Day 1"),
  mutate(z, day = "Day 2")
)
```

```
## # A tibble: 6 x 3
##   x1       x2 day
##   <chr> <int> <chr>
## 1 A         1 Day 1
## 2 B         2 Day 1
## 3 C         3 Day 1
```

```
## 4 B         2 Day 2
## 5 C         3 Day 2
## 6 D         4 Day 2
```

3. The rows contained in `z` are bad! Remove those rows from `y`. set diff or anti join work here setdiff(y removes from z =(y,z))

```
setdiff(y, z)
```

```
## # A tibble: 1 x 2
##   x1       x2
##   <chr> <int>
## 1 A         1
```